

Factors of Box Office Success

A STA610 Case Study

Andrew Amore

2022-10-14

Abstract

Movie production companies are interested in understanding what variables contribute to a successful film, particularly whether a film's budget and IMBD score are predictive of net profits. Box office data from all 2019 film releases was collected and enhanced with additional fields to address this question.

A linear model, with no random effects, was initially fit to explore response associations and establish a baseline. The initial model displayed poor fit according to residual standard error and additional covariate inclusion/interaction was explored, warranting a more complicated analysis. Several hierarchical models were assessed to address heteroscedastic variance concerns using different data grouping structures. Ultimately, random effects were applied to observations based on release date (month). The random intercept model showed little estimation difference compared to the reference model, but adding a random slope on 'Critic.Score' enhanced the estimation accuracy for most fixed effects.

In the final model estimates for 'Budget' and 'Critic.Score' have significant positive effects on profitability at the 95% level. Every \$1 increase in budget yields an additional profit of \$2.43 and improving the IMBD score by 1 point leads to a net increase in profit of ~\$45 million. In addition, 'Title.Sentiment', defined as the polarity score of the film's name, had a surprising p-value of 0.057 in the final model indicating a ~\$97 million net profit increase for every additional polarity score increment. Something a production company should consider when naming a movie!

Dataset Overview & Issues

Metadata Information

All fields used in the analysis are shown in Table 1 with metadata descriptions and a sample observation. Fields highlighted in yellow are augmentations.

Table 1: Metadata Information

Field Name	Description	Sample
Title	film's title	<i>Isn't It Romance</i>
Release.Date	date of movie release	2019-02-13
Production.Company	name of the movie production company	Warner Bros. Pictures / New Line Cinema / Netflix / Bron Creative
Lead.Cast.1	leading actor/actress in the film	Rebel Wilson
Lead.Cast.2	second leading actor/actress in the film	Liam Hemsworth
Lead.Cast.3	third leading actor/actress in the film	Adam DeVine
Director	director of the film	Todd Strauss-Schulson
Box.Office	amount (USD) film made at the box office	48800000
Budget	budget (USD) of film	31000000
Run.Time	run time (minutes) of film	89
Critic.Score	IDMB score of film	5.9
Genre	full genre	Romance/Comedy
Genre.1	main genre of film	Romance
Genre.2	secondary genre of film	Comedy
Num.Genres	number of genres of film	2
Net.Profit	profit (USD) of film	17800000
Title.Sentiment	polarity score of movie title	-0.289
Day.of.Week	release day	Wednesday
Release.Month	release month	February

Missing Data

Exploratory data analysis (EDA) revealed a large number of missing data across a subset of fields. Figure 1 shows the five fields with missing data. Over a third of movie releases are missing “Net.Profit”...the response variable we want to investigate! Clearly something must be done to address this issue. Bayesian imputation methods were investigated to “draw” missing data from candidate probability distributions derived from available data, but due to the large proportion of missing data in the response variable this strategy was abandoned in favor of dropping affected observations. This avoids adding any statistical bias from an imputation method, however, this increases variability of results. Before proceeding several additional data quality issues were discovered.




Figure 1: Missing Data Summary

Limited Releases: There are some additional concerns with the data. First, we have no information about how box office metrics were collected, such as how long after a film’s release date or the number of screening theaters. Films debuting as limited releases have fewer opportunities to generate box office profits and have skewed revenue numbers. A known example of a limited release is “The Irishman”, shown below. At first glance the film appears to have lost over \$100 million, but additional research indicates the film only released at eight theaters nationally. Because of this concern, all Netflix films were dropped from analysis except “Isn’t it Romance” as secondary research confirmed a full theatrical release.

Table 2: Outlier Sample

Title	Production.Company	Box.Office	Budget	Net.Profit	Removal.Reason
Avengers: Endgame	Marvel Studios	2,798,000,000	365,000,000	2,433,000,000	Outlier
The Lion King	Walt Disney Pictures	1,663,000,000	250,000,000	1,413,000,000	Outlier
The Irishman	Netflix / TriBeCa Productions	8,000,000	159,000,000	-151,000,000	Limited Release
Frozen II	Walt Disney Pictures / Walt Disney Animation Studios	1,450,000,000	150,000,000	1,300,000,000	Outlier

Outliers: The dataset also contains a number of outliers which can lead to large leverage points and/or violate modeling assumptions. One such example is “Avengers: Endgame”, which released as the final installment to a series of films from the past decade to much fanfare and the box office numbers are almost twice as large as the next closest film. It’s reasonable to assume that many people were motivated by a different set of factors to watch this film and as a result it should be included in the analysis. Other “sequels” with extraneous factors were also removed (Lion King & Frozen II) for the same reason.

Variable Selection

The main objective of this analysis is to investigate factors influencing ‘Net.Profit’. Before beginning to build a model the covariance structure of the dataset was investigated for films with no missing values. Table 3 displays the

correlation coefficients for numeric fields and indicates mostly weak relationships (correlation < 0.5) with the response. Unsurprisingly, 'Box.Office' and 'Budget' show the strongest relationship. Figure 2 shows the relationship between the response and the two main covariates of interest: budget and critical score.

Table 3: Correlation Matrix

	Box.Office	Budget	Run.Time	Critic.Score	Num.Genres	Net.Profit	Title.Sentiment
Box.Office	1	0.77	0.27	0.3	-0.04	0.99	0.06
Budget	0.77	1	0.28	0.14	-0.05	0.65	-0.09
Run.Time	0.27	0.28	1	0.47	-0.06	0.25	0.07
Critic.Score	0.3	0.14	0.47	1	0.03	0.32	0.18
Num.Genres	-0.04	-0.05	-0.06	0.03	1	-0.03	-0.04
Net.Profit	0.99	0.65	0.25	0.32	-0.03	1	0.1
Title.Sentiment	0.06	-0.09	0.07	0.18	-0.04	0.1	1




Figure 2: Relationship with Net-Profit

Model Building

Standard Regression

Before more complex modeling, a baseline linear regression model was fit on the data. Let $y \in \mathbb{R}^n$ be a vector of movie net profits, $X \in \mathbb{R}^{n \times 2}$ be a matrix of covariates ('Budget' & 'Critic.Score'), and β_0, β be regression coefficients and intercept respectively. The model can be formalized as follows:

$$y = \beta_0 + X\beta + \epsilon, \text{ where } \epsilon_{iid} \sim N(0, \sigma^2 I)$$

The results of the model fit including coefficient estimates, uncertainty quantification via 95% confidence interval and p-values are shown in Table 4. All coefficients, including the intercept are shown to be statistically significant, but standard errors (SE) are large for Critic.Score and intercept.

Table 4: Initial Results

Coefficient	Estimate (USD)	SE (USD)	P-Value	CI Lower (2.5%)	CI Upper (97.5%)
(Intercept)	-315,071,394	82,493,418	0.000197	-478,106,904	-152,035,885
Budget	2.38	0.233	7.71e-19	1.92	2.84
Critic.Score	49,072,204	12,985,672	0.000229	23,408,027	74,736,381

Observations	Residual Std. Error	R ²	Adjusted R ²
149	147,033,657	0.476	0.469

Based on the high SEs, the model fit was assessed. The scale location plot is the most interesting as it addresses the large differences in film net profits which were left as raw dollar amounts. Clearly, as the fitted value increases the standardized residual increases which is a violation of the linear model homoscedastic variance assumption. Before moving to a more complex modeling technique additional covariates were added to try and improve the baseline regression.




Figure 3: Standardized Residual Plot

Adding Additional Covariates to the Standard Regression

The correlation matrix in Table 3 suggests ‘Run.Time’ and ‘Title.Sentiment’ may be good predictors to add to the model. Models with each additional covariate were compared to the reference model using an F-test to determine inclusion criteria. In addition to the two numeric covariates, ‘Day.of.Week’ and ‘Genre.1’ were also investigated. F-test summaries are provided below and based on these results only ‘Title.Sentiment’ was added to the model as closer inspection of ‘Day.of.Week’ revealed a high variance estimate from Tuesday (only has three observations).

Table 5: F-Test Results

Added Covariate	Df	RSS	RSS Reduction from Ref.	Test-Statistic (F)	P-Value
Reference Model	146	3.16e+18	NA	NA	NA
Title.Sentiment	145	3.08e+18	7.21e+16	3.390	0.0676
Day.of.Week	143	2.99e+18	1.68e+17	2.683	0.0491
Run.Time	145	3.15e+18	1.08e+16	0.496	0.4824
Genre.1	117	2.31e+18	8.50e+17	1.487	0.0725

The final “baseline” linear regression model results are shown in Table 6.

Table 6: Final Baseline Reference Results

Coefficient	Estimate (USD)	SE (USD)	P-Value	CI Lower (2.5%)	CI Upper (97.5%)
(Intercept)	-282,789,925	83,683,651	0.000934	-448,187,277	-117,392,573
Budget	2.43	0.232	2.09e-19	1.97	2.89
Critic.Score	44,256,400	13,143,548	0.000973	18,278,709	70,234,091
Title.Sentiment	93,913,069	51,006,166	0.0676	-6,898,556	194,724,694

Observations	Residual Std. Error	R ²	Adjusted R ²
149	145,844,747	0.488	0.478

This final model improves the residual standard error estimate from the reference model, but is still a relatively poor fit (high coefficient SEs). To try and reduce these values further hierarchical modeling is explored.

Hierarchical Modeling

To perform an adequate hierarchical analysis a “clustering” category is needed. At first glance, ‘Genre’ shows promise, but many films are recorded with multiple or unique even entries with limited overlap. Table 7 displays the number of films with multiple genres. The median film has more than one genre, making grouping more challenging.

Table 7: Genre Clustering

Number of Genres by Film	Number of Observations
1	81
2	110
3	17
4	1

One could circumvent this by using the first listed entry as the primary genre and proceed with analysis, but there are still quite a few clusters with few observational data points shown in Figure 4. A proposed solution is to categorize similar genres into a hierarchy for analysis. For example, maybe “War” and “Action” involve enough similarity to warrant combining into one super-category, however, extensive domain knowledge may be required to infer less obvious groupings. This procedure might introduce statistical bias that could alter results.




Figure 4: Count of Films by Genre.1

A more objective hierarchy is needed and one hypothesis is that audiences may be sensitive to the seasonality of a film’s release. For example, during the winter months individuals might be more likely to attend a new film because there are less alternative outdoor activities compared to other seasons. Production companies might alter a film’s budget accordingly. Figure 5 displays linear regression fits contrasting the group effect. Clearly, there are differences in Net.Profit trends across groups, warranting the hierarchical model.

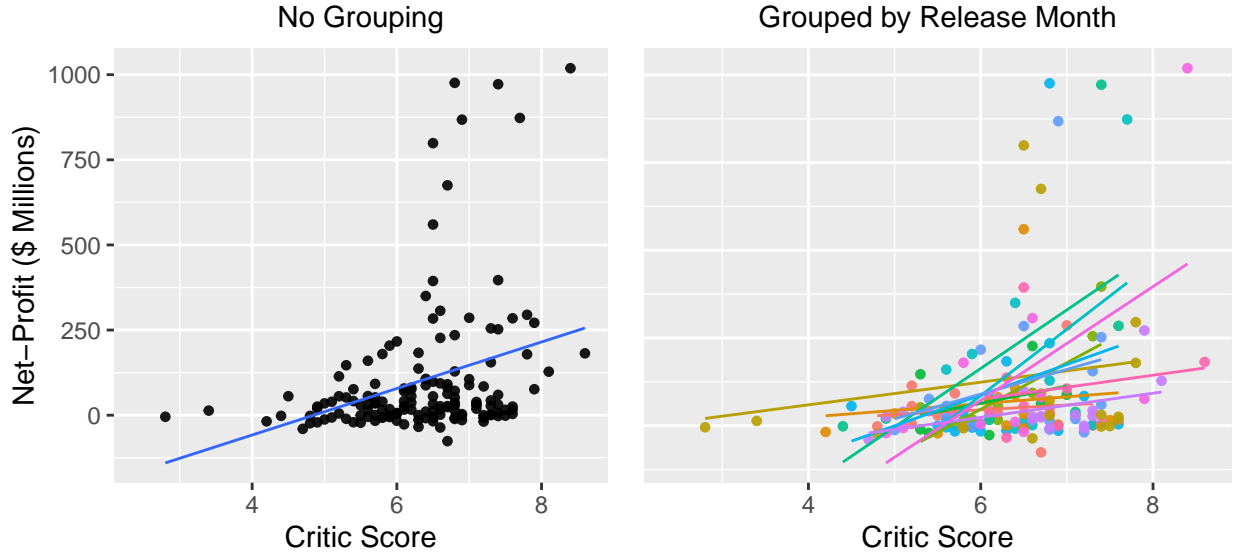


Figure 5: Motivating the Hierarchical Model

Several hierarchical models were tried with the data. First, a random intercept on release month was added. The model can be described using the same notation as the final reference linear model with new parameter $\alpha_j \in \mathbb{R}^{12 \times 1}$ denoting the random effect for release month j .

$$y = \beta_0 + \alpha_j + X\beta + \epsilon, \text{ where } \epsilon_{iid} \sim N(0, \sigma^2 I)$$

This model detects very little differences with the addition of the random effect and all parameter estimates are almost identical to the reference linear model.

Table 8: Mixed Model 1 Results

Coefficient	Estimate (USD)	SE (USD)	P-Value	CI Lower (2.5%)	CI Upper (97.5%)
(Intercept)	-282,789,925	83,683,651	0.000727	-445,638,767	-1.2e+08
Budget	2.43	0.232	1.45e-25	1.98	2.88
Critic.Score	44,256,400	13,143,548	0.000759	18,678,984	69,833,816
Title.Sentiment	93,913,069	51,006,166	0.0656	-5,345,209	193,171,348

Coefficient	CI Lower (2.5%)	CI Upper (97.5%)
sd_(Intercept) Release.Month	0	Inf
sigma	128,969,424	161,893,468

Clearly the model doesn’t detect any differences under the first model and ideally, we would like to add random slope terms for all covariates, but unfortunately, the small sample size can lead to boundary fits and singularity for too many parameters.

The final hierarchical model is defined with similar notation as follows. To avoid boundary conditions, the random intercept is removed and instead estimate the random slope for ‘Critic.Score’, $\gamma_{j[i]}$ the random effect for critic score of film c_i , in addition to all other parameters from the baseline reference linear model.

$$y_{ij} = \beta_0 + X\beta + \gamma_{j[i]}c_i + \epsilon_i \text{ where } \epsilon_i \sim N(0, \sigma^2)$$

Accounting for this new parameter alters the parameter estimates compared to the baseline reference model, but does not dramatically change the coefficient SE estimates... a disappointment. Comparing Table 9 to Table 6 the estimate for 'Budget' is identical, but 'Critic.Score' and 'Title.Sentiment' see an increase with lower p-values that almost nudge 'Title.Sentiment' into a statistically significant effect at the 95% level.

Table 9: Mixed Model 2 Results

Coefficient	Estimate (USD)	SE (USD)	P-Value	CI Lower (2.5%)	CI Upper (97.5%)
(Intercept)	-287,034,469	83,429,515	0.000584	-448,679,088	-121,223,595
Budget	2.43	0.232	1.23e-25	1.98	2.88
Critic.Score	45,088,054	13,131,958	0.000605	18,880,324	70,552,765
Title.Sentiment	96,888,305	50,903,239	0.0571	-4,563,154	195,533,850

Coefficient	CI Lower (2.5%)	CI Upper (97.5%)
sd_Critic.Score Release.Month	0	8,369,437
sigma	127,953,194	161,739,097

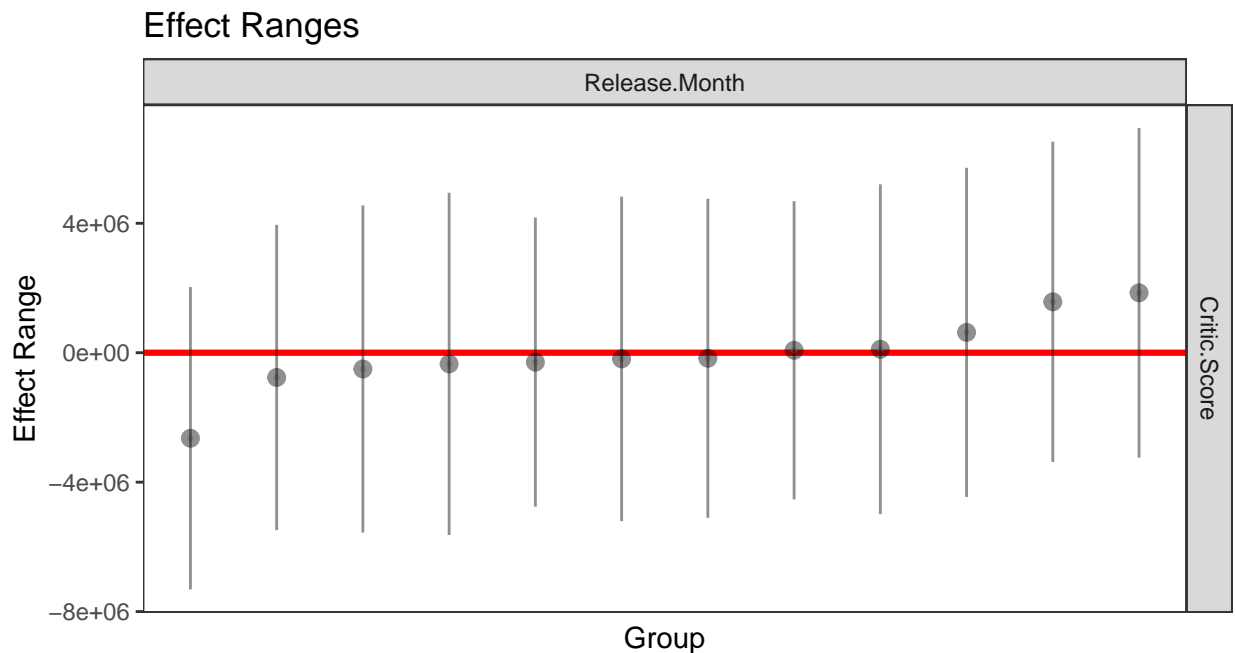


Figure 6: Interval Estimates for the Random Effect

Conclusions, Limitations & Enhancements

In this analysis factors affecting box office success were evaluated. Before modeling, covariates were cleaned, analyzed and enhanced to better understand dependence. Roughly 100 observations were dropped from the analysis related to data quality. Standard linear regression and mixed effect models were evaluated to assess the statistical significance of covariates on profitability. 'Budget' and 'Critic.Score' were found to have significant positive effects at the 95% confidence level, while 'Title.Sentiment' was found to be significant at the 90% confidence level also adding a positive effect. Obviously, it is not always possible to create a high budget, critically acclaimed film, but one actionable recommendation to production companies is to title films more positively as audiences seem more inclined to watch these films. However, it's important to note that more analysis is needed to determine the causal relationship between 'Title.Sentiment' and net profitability.

There are several limitations of this analysis mentioned throughout, but the main concerns are highlighted. First, the sample contains a large fraction of missing data, is relatively small, and only contains a single year of observations

which increases the variability of the analysis as 2019 might not be representative year of overall film releases. Second, no knowledge of the data collection procedure is provided and increases the difficulty of assessing covariate effects on overall net profit because some films might release in a limited number of theaters or for a pre-specified length of time that makes aggregate profit potentially misleading. To combat this one could code profitability as a binary response indicator, $\{0, 1\}$, and fit a logistic regression on the new dependent variable. This would minimize the influence of large outliers observed in the dataset and potentially make the analysis more robust. However, one would lose the ability to quantify the direct covariate effect on overall profit which feels relevant to evaluate. Lastly, there is no obvious clustering field for a hierarchical model provided in the dataset and spending more time to group genres together rather than month of release could yield a more beneficial analysis.